

2 | IMPORT NECESSARY LIBRARIES

In [1]:

```

import pandas as pd  #for data manipulation operations
import numpy as np  #for numeric operations on data
import seaborn as sns  #for data visualization operations
import matplotlib.pyplot as plt  #for data visualization operations
from sklearn.preprocessing import LabelEncoder # for encoding
from sklearn.preprocessing import StandardScaler #for standardization

from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.metrics import plot_confusion_matrix
from sklearn import model_selection
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from termcolor import colored
#!pip install xgboost
from xgboost import XGBRegressor

#!pip install lightgbm
from lightgbm import LGBMRegressor

#ignore warnings
import warnings
warnings.filterwarnings("ignore")

from sklearn import set_config
set_config(print_changed_only = False)

sns.set_theme(style = "whitegrid")

print(colored("\n THE REQUIRED LIBRARIES WERE SUCCESFULLY IMPORTED...", color = "green", attrs = ["bold", "dark"]))

```

THE REQUIRED LIBRARIES WERE SUCCESFULLY IMPORTED...

3|LOAD DATASET

In [2]:

```
heart = pd.read_csv("../input/heart-failure-clinical-data/heart_failure_c  
linical_records_dataset.csv")  
df = heart.copy()  
df.head(n = 10).style.background_gradient(cmap = "Reds_r").set_properties  
(**{"font-family" : "Segoe UI"}).hide_index()
```

Out[2]:

age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure
75.000000	0	582	0	20	1
55.000000	0	7861	0	38	0
65.000000	0	146	0	20	0
50.000000	1	111	0	20	0
65.000000	1	160	1	20	0
90.000000	1	47	0	40	1
75.000000	1	246	0	15	0
60.000000	1	315	1	60	0
65.000000	0	157	0	65	0
80.000000	1	123	0	35	1

4 | INITIAL INFORMATION ABOUT DATASET

4.1 | Get Initial Information

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 299 entries, 0 to 298
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   299 non-null    float64
1   anaemia                              299 non-null    int64
2   creatinine_phosphokinase             299 non-null    int64
3   diabetes                             299 non-null    int64
4   ejection_fraction                    299 non-null    int64
5   high_blood_pressure                  299 non-null    int64
6   platelets                            299 non-null    float64
7   serum_creatinine                     299 non-null    float64
8   serum_sodium                         299 non-null    int64
9   sex                                  299 non-null    int64
10  smoking                              299 non-null    int64
11  time                                 299 non-null    int64
12  DEATH_EVENT                          299 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 30.5 KB
```

Brief information

The dataset consists of **299 rows and 13 columns**. The data type of all variables are **numeric**.

4.2 | Change Column Values

In [4]:

```
df_eda = pd.DataFrame()

df_eda["age"] = df["age"]
df_eda["anaemia"] = np.where(df["anaemia"] < 1, "no", "yes")
df_eda["creatinine_phosphokinase"] = df["creatinine_phosphokinase"]
df_eda["diabetes"] = np.where(df["diabetes"] < 1, "no", "yes")
df_eda["ejection_fraction"] = df["ejection_fraction"]
df_eda["high_blood_pressure"] = np.where(df["high_blood_pressure"] < 1,
"no", "yes")
df_eda["platelets"] = df["platelets"]
df_eda["serum_creatinine"] = df["serum_creatinine"]
df_eda["serum_sodium"] = df["serum_sodium"]
df_eda["sex"] = np.where(df["sex"] < 1, "female", "male")
df_eda["smoking"] = np.where(df["smoking"] < 1, "no", "yes")
df_eda["time"] = df["time"]
df_eda["death_event"] = np.where(df["DEATH_EVENT"] < 1, "no", "yes")

df_eda.head().style.background_gradient(cmap = "Reds").set_properties(**
{"font-family" : "Segoe UI"}).hide_index()
```

Out[4]:

age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure
75.000000	no	582	no	20	yes
55.000000	no	7861	no	38	no
65.000000	no	146	no	20	no
50.000000	yes	111	no	20	no
65.000000	yes	160	yes	20	no

4.3 | Descriptive Statistics of Numeric Variables

```
In [5]: df_eda.describe().T.style.background_gradient(cmap = "Reds_r").set_properties(**{"font-family" : "Segoe UI"})
```

Out[5]:

	count	mean	std	min	25%
age	299.000000	60.833893	11.894809	40.000000	51.000000
creatinine_phosphokinase	299.000000	581.839465	970.287881	23.000000	116.500000
ejection_fraction	299.000000	38.083612	11.834841	14.000000	30.000000
platelets	299.000000	263358.029264	97804.236869	25100.000000	212500.000000
serum_creatinine	299.000000	1.393880	1.034510	0.500000	0.900000
serum_sodium	299.000000	136.625418	4.412477	113.000000	134.000000
time	299.000000	130.260870	77.614208	4.000000	73.000000

Basic descriptive statistics

- The average value of age is **60.83** , the highest value is **95**
- The average value of creatinine_phosphokinase is **581.83** , the highest value is **7861**
- The average value of ejection_fraction is **30.08** , the highest value is **80**
- The average value of platelets is **263358** , the highest value is **850000**
- The average value of serum_creatinine is **1.39** , the highest value is **9.4**
- The average value of serum_sodium is **136.62** , the highest value is **148**
- The average value of time is **130.26** , the highest value is **285**

4.4 | Check null Values

```
In [6]: df.isnull().any() #to check "null" values
```

```
Out[6]:
```

age	False
anaemia	False
creatinine_phosphokinase	False
diabetes	False
ejection_fraction	False
high_blood_pressure	False
platelets	False
serum_creatinine	False
serum_sodium	False
sex	False
smoking	False
time	False
DEATH_EVENT	False

dtype: bool

5| DATA VISUALIZATION

In [15]:

```
plt.figure(figsize = (18, 10), dpi = 50)

sns.barplot(x = "anaemia", y = "creatinine_phosphokinase", hue = "sex",
            data = df_eda, palette = "Set1", saturation = 1);
```




Chart report

From the graphs above, we can observe the relationships between various categorical and numerical variables.

5.2 | Distplot

In [16]:

```
fig, axes = plt.subplots(1, 2, figsize=(20, 7))

sns.distplot(ax = axes[0], x = df_eda["age"],
             hist = True,
             bins = 20,
             kde = True,
             vertical = False, color = "#C44D04").set(title = "The distribution of the values of 'age' variable");

sns.distplot(ax = axes[1], x = df_eda["platelets"],
             hist = True,
             bins = 20,
             kde = True,
             vertical = False, color = "#FA8B47").set(title = "The distribution of the values of 'platelets' variable");
```

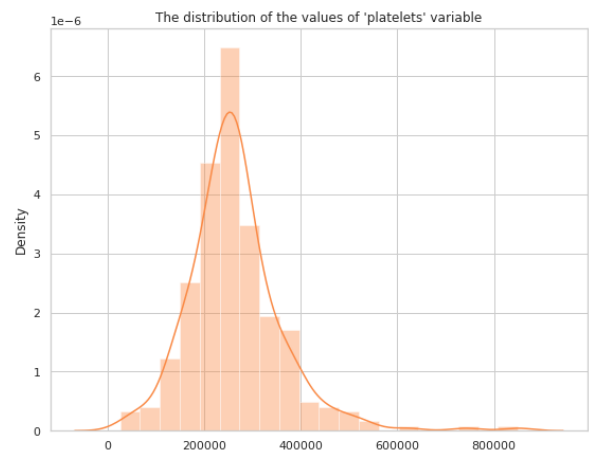





Chart report

From the first graph, we can observe that the values of the "age" variable are mostly distributed as values of 50, 60 and 65 - 70. In addition, there are many peaks in the distribution graph of the "age" variable. This is an indicator that the values of the variable are not normally distributed. Looking at the second graph, we can observe that the values of the "platelets" variable are mainly distributed between 200000 and 400000. At the same time, the graph does not have many peaks, skewness and kurtosis are low. It can be said that the values of this variable are normally distributed.

In [17]:

```
sns.displot(
    data = df_eda,
    x = "age", hue = "sex",
    kind="kde", height=6,
    multiple="fill", clip=(0, None),
    color = "#FE6203",
).set(title = "density of the classes of 'sex' variable by age");

sns.displot(
    data = df_eda,
    x = "age", hue = "death_event",
    kind = "kde", height=6,
    multiple="fill", clip=(0, None),
    palette="ch:rot=-.25,hue=1,light=.75",
).set(title = "density of the classes of 'death_event' variable by age");
```




5.3 | Boxplot

In [18]:

```
fig, axes = plt.subplots(1, 2, figsize=(20, 10))

sns.boxplot(ax = axes[0], x = "sex", y = "serum_sodium", data = df_eda, width = 0.7,
            orient = "v", fliersize = 3, linewidth = 1);
axes[0].set_title("relationship between 'sex' and 'serum-sodium' variables");


sns.boxplot(ax = axes[1], x = "sex", y = "platelets", data = df_eda, width = 0.7,
            orient = "v", fliersize = 3, linewidth = 1);
axes[1].set_title("relationship between 'sex' and 'platelets' variables");
```



GG.plot

In [33]:

```
plt.figure(figsize = (18, 10), dpi = 50)
plt.style.use('ggplot')
sns.lineplot(data = df_eda, x = "age", y = "time", hue = "sex" );
```



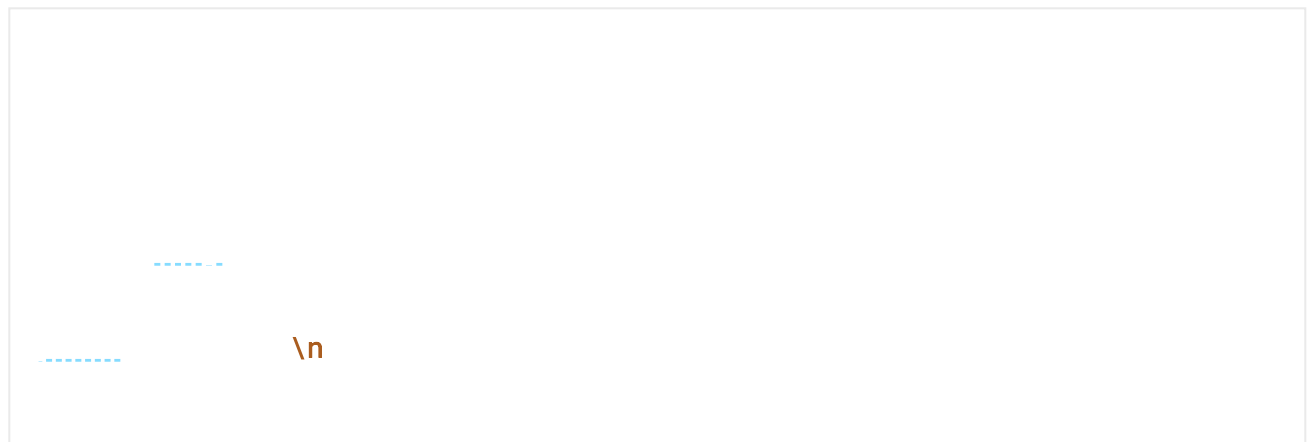
6.1 | Look at Dataset

In [34]:

```
df.head().style.background_gradient(cmap = "Reds_r")
```

Out[34]:

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_press
0	75.000000	0	582	0	20	1
1	55.000000	0	7861	0	38	0
2	65.000000	0	146	0	20	0
3	50.000000	1	111	0	20	0
4	65.000000	1	160	1	20	0



DEPENDENT AND INDEPENDENT VARIABLES WERE SUCCESFULLY SELECTED...

6.3 | Split Dataset into Train and Test Sets

In [36]:

```
x_train, x_test, y_train, y_test = train_test_split(x, y,
                                                    test_size = 0.20,
                                                    shuffle = True,
                                                    random_state = 1)

print(colored("\n THE DATASET WERE SUCCESFULLY SPLITTED - test = 20%, tra
in = 80%...", color = "green", attrs = ["bold", "dark"]))
```

THE DATASET WERE SUCCESFULLY SPLITTED - test = 20%, train = 80%...

6.4 | Standardization

In [37]:

```
scaler = StandardScaler()
scaler.fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)

print(colored("\n x_train AND x_test SETS WERE SUCCESFULLY STANDARIZED",
color = "green", attrs = ["bold", "dark"]))
```

x_train AND x_test SETS WERE SUCCESFULLY STANDARIZED

In [38]:

```
x_train[0:5]
```

Out[38]:

```
array([[ -1.33194278,  1.11069566, -0.4576634 , -0.84818893, -0.2431805
8,
        -0.72269841, -0.89319445, -0.58084862,  0.51465589,  0.7294184
5,
        1.52297224, -0.10575054],
 [-0.4999137 , -0.90033664, -0.51894319, -0.84818893,  0.1773704
9,
        -0.72269841, -0.57954879, -0.39670905,  0.29500811,  0.7294184
5,
        -0.65661079,  1.33189826],
 [ 0.58172409,  1.11069566,  0.39546617,  1.17898261, -0.2431805
8,
        -0.72269841,  0.15589621, -0.30463927, -0.58358301,  0.7294184
5,
        -0.65661079,  0.86979686],
 [-0.08389917,  1.11069566, -0.02487488, -0.84818893, -0.6637316
5,
        1.38370306, -1.40151673, -0.48877884,  1.83254257, -1.3709551
7,
        -0.65661079, -0.43949044],
 [-0.08389917, -0.90033664, -0.35712624,  1.17898261,  0.0091500
7,
        -0.72269841,  0.78318753,  1.44468663,  1.17359923, -1.3709551
7,
        -0.65661079, -1.27384019]])
```


In [39]:

```
x_test[0:5]
```

Out[39]:


```
array([[ -0.91592824,  1.11069566, -0.47202585, -0.84818893, -1.5048338
,
        -0.72269841, -0.73096394, -0.58084862,  0.51465589,  0.7294184
5,
        -0.65661079,  0.21515321],
 [-1.33194278, -0.90033664, -0.02487488,  1.17898261,  1.4390237
1,
        -0.72269841,  3.09767621, -0.39670905, -1.02287857, -1.3709551
7,
        -0.65661079,  1.55011281],
 [-0.66631952,  1.11069566, -0.49500577, -0.84818893, -1.5048338
,
        1.38370306,  1.74575525, -0.02842992,  0.51465589, -1.3709551
7,
        -0.65661079, -1.10697024],
 [-0.74952242, -0.90033664, -0.4557484 , -0.84818893, -0.6637316
5,
        -0.72269841, -0.41731828, -0.67291841, -0.14428745,  0.7294184
5,
        1.52297224, -0.22127589],
 [ 0.33211537,  1.11069566, -0.45287591, -0.84818893, -0.2431805
8,
        1.38370306,  0.36138819, -0.58084862, -0.58358301,  0.7294184
5,
        -0.65661079,  0.83128841]])
```

In [48]:

```
# Visualization Confusion Matrix
conf_mat = confusion_matrix(y_test, y_pred)
print(conf_mat)


# Visualize it as a heatmap
import seaborn
seaborn.heatmap(conf_mat, square = True, robust = True)
plt.show()
```

```
[[42  4]
 [ 5  9]]
```



In [49]:

```
plot_confusion_matrix(rf_model,  
                      x_test,  
                      y_test,  
                      cmap = plt.cm.Reds,  
                      normalize='true');
```



7.10 | Logistic Regression



What is logistic regression?

This type of statistical model (also known as logit model) is often used for classification and predictive analytics. Logistic regression estimates the probability of an event occurring, such as voted or didn't vote, based on a given dataset of independent variables. Since the outcome is a probability, **the dependent variable is bounded between 0 and 1**. In logistic regression, a logit transformation is applied on the odds - that is, the probability of success divided by the probability of failure. This is also commonly known as the log odds, or the natural logarithm of odds, and this logistic function is represented by the following formulas:

$$\text{Logit}(\pi) = 1/(1 + \exp(-\pi))$$

$$\ln(\pi/(1-\pi)) = \text{Beta}_0 + \text{Beta}_1 X_1 + \dots + \text{Beta}_k X_k$$

In this logistic regression equation, **logit(π) is the dependent or response variable and x is the independent variable**. The beta parameter, or coefficient, in this model is commonly estimated via maximum likelihood estimation (MLE). This method tests different values of beta through multiple iterations to optimize for the best fit of log odds. All of these iterations produce the log likelihood function, and logistic regression seeks to maximize this function to find the best parameter estimate. **Once the optimal coefficient (or coefficients if there is more than one independent variable) is found, the conditional probabilities for each observation can be calculated, logged, and summed together to yield a predicted probability**. For binary classification, a probability less than .5 will predict 0 while a probability **greater than 0 will predict 1**. After the model has been computed, it's best practice to evaluate the how well the model predicts the dependent variable, which is called goodness of fit. The Hosmer-Lemeshow test is a popular method to assess model fit.



In [50]:

```
# build a model

log = LogisticRegression(solver = "liblinear")
log_model = log.fit(x_train, y_train)
log_model
```

Out[50]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept
= True,

                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001,
                    verbose=0,

                    warm_start=False)
```

7.11 | Accuracy Score of Model on Test Set

In [51]:

```
y_pred = log_model.predict(x_test)
accuracy_score(y_test, y_pred)
```

Out[51]:

```
0.8833333333333333
```

7.12 | ROC AUC - Logistic Regression

In [52]:

```
log_roc_auc = roc_auc_score(y_test, log_model.predict(x_test))

fpr, tpr, thresholds = roc_curve(y_test, log_model.predict_proba(x_test)
[:,1])
plt.figure()
plt.plot(fpr, tpr, label = 'AUC (area = %0.2f)' % log_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
plt.show()
```

